

Optimization Models in Machine Learning: Introduction and Examples

University of Washington

July 30, 2020

Outline

1. Intro to Optimization + Notebook 1
2. Linear Regression and Regularization + Notebook 2
3. Logistic Regression + Notebook 3
4. Outlier Removal + Notebook 4

Optimization: Overview

A general optimization problem has the form

$$\begin{array}{ll}\underset{x}{\text{minimize}} & f_0(x) \\ \text{subject to} & f_i(x) \leq b_i, \quad i = 1, \dots, m,\end{array}$$

with components

- ▶ $x = (x_1, \dots, x_n)$ - optimization variable
- ▶ $f_0 : \mathbf{R}^n \rightarrow \mathbf{R}$ - objective function
- ▶ $f_i : \mathbf{R}^n \rightarrow \mathbf{R}$ - constraint functions; b_i - constraint bounds

Optimization: Overview

A general optimization problem has the form

$$\begin{array}{ll}\underset{x}{\text{minimize}} & f_0(x) \\ \text{subject to} & f_i(x) \leq b_i, \quad i = 1, \dots, m,\end{array}$$

with components

- ▶ $x = (x_1, \dots, x_n)$ - optimization variable
- ▶ $f_0 : \mathbf{R}^n \rightarrow \mathbf{R}$ - objective function
- ▶ $f_i : \mathbf{R}^n \rightarrow \mathbf{R}$ - constraint functions; b_i - constraint bounds

Many applications:

- ▶ Data fitting and regression
- ▶ Classification
- ▶ Image processing
- ▶ Portfolio optimization
- ▶ Recommender systems
- ▶ Optimal control
- ▶ Medical treatment planning

Optimization: Problem Classes

There are different classes of optimization problems, which can determine a problem's difficulty and solution method:

- ▶ Constrained vs. Unconstrained
- ▶ Smooth vs. Nonsmooth
- ▶ Convex vs. Nonconvex

¹Stephen Boyd and Lieven Vandenberghe. Convex Optimization. Cambridge University Press, 2004.

Optimization: Problem Classes

There are different classes of optimization problems, which can determine a problem's difficulty and solution method:

- ▶ Constrained vs. Unconstrained
- ▶ Smooth vs. Nonsmooth
- ▶ Convex vs. Nonconvex

An important class: **convex optimization** problems

“With only a bit of exaggeration, we can say that if you formulate a practical problem as a convex optimization problem, then you have solved the original problem.”¹

¹Stephen Boyd and Lieven Vandenberghe. Convex Optimization. Cambridge University Press, 2004.

Optimization: Problem Classes

A convex optimization problem has objective and constraint functions that satisfy the inequality

$$f_i(\lambda x + (1 - \lambda)y) \leq \lambda f_i(x) + (1 - \lambda)f_i(y)$$



for all $x, y \in \mathbf{R}^n$ and all $0 \leq \lambda \leq 1$.

Optimization: Problem Classes

A convex optimization problem has objective and constraint functions that satisfy the inequality

$$f_i(\lambda x + (1 - \lambda)y) \leq \lambda f_i(x) + (1 - \lambda)f_i(y)$$

for all $x, y \in \mathbf{R}^n$ and all $0 \leq \lambda \leq 1$.



Important consequence: in a convex problem, no “local minima”

Optimization: Solution Methods

Very few optimization problems have a closed-form solution (e.g., least-squares); most problems are solved using iterative methods.

One important iterative method is gradient descent:



$$x^{k+1} = x^k - \alpha \nabla f(x^k)$$

Optimization: Solution Methods

Very few optimization problems have a closed-form solution (e.g., least-squares); most problems are solved using iterative methods.

One important iterative method is gradient descent:

$$x^{k+1} = x^k - \alpha \nabla f(x^k)$$



Optimization: Machine Learning Models

Many problems in machine learning seek to build a model

$$g(a; x) \approx y$$

given a data set

$$\left\{ (a_1, y_1), \dots, (a_m, y_m) \right\},$$

with components

- ▶ $a_i = (a_{i1}, \dots, a_{in})$ - data features
- ▶ $y_i \in \mathbf{R}$ or $\{0, 1\}$ - data value or label/class
- ▶ $g : \mathbf{R}^n \rightarrow \mathbf{R}$ or $\{0, 1\}$ - prediction function
- ▶ $x = (x_1, \dots, x_n)$ - model parameters
- ▶ m - number of data points
- ▶ n - number of data features

Optimization: Machine Learning Models

We can fit a model to the given data by solving an optimization problem of the form

$$\underset{x}{\text{minimize}} \quad \sum_{i=1}^m f_i(g(a_i; x), y_i) + r(x),$$


with components

- ▶ $x = (x_1, \dots, x_n)$ - model parameters we want to learn
- ▶ $f_i : \mathbf{R}^n \rightarrow \mathbf{R}$ - “loss” functions: measure how well the model fits the data for given parameters; e.g., $(g(a_i; x) - y_i)^2$
- ▶ $r(x) : \mathbf{R}^n \rightarrow \mathbf{R}$ - regularization function


Optimization: Machine Learning Models

We focus on two common problems in machine learning:

Linear Regression




Logistic Regression (Classification)




Optimization: Machine Learning Models

We focus on two common problems in machine learning:

Linear Regression



Logistic Regression (Classification)



- ▶ Data: Continuous features $\{a_i\}$ and outputs $\{y_i\}$
- ▶ Goal: Find linear predictor

$$x_0 + x_1 a_i \approx y_i$$

- ▶ Approach: Assume a statistical model for errors and develop a maximum likelihood formulation

Linear Regression: Derivation

Assuming the errors in our data come from a normal distribution,

$$y_i = x_0 + x_1 a_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2) \quad \text{independent,}$$

Linear Regression: Derivation

Assuming the errors in our data come from a normal distribution,

$$y_i = x_0 + x_1 a_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2) \quad \text{independent,}$$

the probability of observation (a_i, y_i) given the parameters is

$$P((a_i, y_i); x_0, x_1) \propto \exp\left(\frac{-(y_i - x_0 - x_1 a_i)^2}{2\sigma^2}\right).$$

Linear Regression: Derivation

Assuming the errors in our data come from a normal distribution,



$$y_i = x_0 + x_1 a_i + \epsilon_i, \quad \epsilon_i \sim N(0, \sigma^2) \quad \text{independent,}$$

the probability of observation (a_i, y_i) given the parameters is

$$P((a_i, y_i); x_0, x_1) \propto \exp\left(\frac{-(y_i - x_0 - x_1 a_i)^2}{2\sigma^2}\right).$$


We can fit model parameters by maximizing the likelihood (minimizing the negative log-likelihood):

$$-\log \prod_{i=1}^m \exp\left(-\frac{(y_i - x_0 - x_1 a_i)^2}{2\sigma^2}\right) \propto \sum_{i=1}^m (y_i - x_0 - x_1 a_i)^2$$



Linear Regression: Intuition and Properties

$$\min_{x_0, x_1} \sum_{i=1}^m (y_i - x_0 - x_1 a_i)^2$$



- ▶ Minimize the least-squares distance between observations y_i and predictions $x_0 + x_1 a_i$.
- ▶ The problem is convex, smooth, and easy to solve.
- ▶ Linear regression actually has a closed-form solution, but it is often found more efficiently by iterative algorithms

Regularization: Overview

Many problems in machine learning add a regularization term $r(x)$ to the objective function to

- ▶ incorporate prior knowledge about *structure* in x , e.g., sparsity or smoothness
- ▶ help avoid overfitting,
- ▶ get more robust (to data perturbations) solutions, or
- ▶ improve the stability of the solution process.

Two popular forms of regularized linear regression:



- ▶ Lasso - $\min_x f(x) + \lambda \|x\|_1$, where $\|x\|_1 = \sum_{i=1}^n |x_i|$
- ▶ Ridge - $\min_x f(x) + \lambda \|x\|_2^2$, where $\|x\|_2^2 = \sum_{i=1}^n x_i^2$

Regularization: Geometric Interpretation

Consider the constrained least-squares problem

$$\begin{aligned} & \underset{x}{\text{minimize}} && \frac{1}{2} \|Ax - y\|_2^2 \\ & \text{subject to} && \|x\|_p \leq t. \end{aligned}$$

Choice of norm influences properties of solution x : with $p = 1$, solutions tend to occur on the vertices, where many $x_i = 0$.



Regularization: Relaxed Constraints

We can move the norm from a constraint into the objective function to get

$$\underset{x}{\text{minimize}} \quad \frac{1}{2} \|Ax - y\|_2^2 + \lambda \|x\|_p,$$

where regularization parameter λ balances model error with how much we regularize.

The Lasso ($p = 1$) is often used to find sparse solutions. Ridge regression ($p = 2$) is often used for ill-conditioned problems.

More generally: regularizers can promote other structures: For example, if the parameters form a matrix X , a low-rank matrix is often desired (e.g., the ‘matrix completion problem’ for recommender systems).

Logistic Regression: Overview

- ▶ Data: Continuous features $\{a_i\}$ and discrete labels $y_i \in \{0, 1\}$
- ▶ Goal: Find linear predictor

$$x_0 + x_1 a_i = \begin{cases} \text{positive} & \Rightarrow y_i = 1 \\ \text{negative} & \Rightarrow y_i = 0 \end{cases}$$

- ▶ Approach: Combine Bernoulli model with a linear predictor
- ▶ Examples: Hours studied vs. Pass/Fail, measurements vs. disease

Logistic Regression: Derivation

Rewriting the Bernoulli model in standard form,

$$\begin{aligned}P\left((a_i, y_i); p_i\right) &= p_i^{y_i} (1 - p_i)^{1-y_i} \\&= \exp\left(y_i \log\left(\frac{p_i}{1 - p_i}\right) + \log(1 - p_i)\right),\end{aligned}$$

we can model the term multiplying y_i using our linear predictor,

$$\log\left(\frac{p_i}{1 - p_i}\right) = x_0 + x_1 a_i,$$

which gives us,

$$\log(1 - p_i) = -\log(1 + \exp(x_0 + x_1 a_i)).$$

Combining the above expressions results in the likelihood function

$$\mathcal{L}(x_0, x_1; (a, y)) = \prod_{i=1}^m \exp\left(y_i(x_0 + x_1 a_i) - \log(1 + \exp(x_0 + x_1 a_i))\right).$$

Logistic Regression: Derivation

We can fit our model parameters to the given data by maximizing the likelihood, or by minimizing the negative log-likelihood:

$$-\log \mathcal{L}(x_0, x_1; (a, y)) = \sum_{i=1}^m \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$$


Explicitly, we solve the following problem

$$\min_{x_0, x_1} \sum_{i=1}^m \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$$

Logistic Regression: Intuition and Properties


$$\min_{x_0, x_1} \sum_{i=1}^m \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$$

- ▶ If the label is 0, we want to make $\log(1 + \exp(x_0 + x_1 a_i))$ as small as possible, equivalent to making $x_0 + x_1 a_i \ll 0$
- ▶ If the label is 1, can show objective decreases with respect to $x_0 + x_1 a_i$, so we want $x_0 + x_1 a_i \gg 0$



Logistic Regression: Intuition and Properties

- We look for intercept x_0 and slope x_1 that do the best job for all the data in the set.



Logistic Regression: Intuition and Properties

- ▶ The problem is convex and smooth, and 'nice' to solve.
- ▶ For a future data points with feature a , $p = \frac{\exp(x_0 + x_1 a)}{1 + \exp(x_0 + x_1 a)}$
- ▶ Other methods can also be used, e.g. support vector machines.

Trimming: Outlier Removal

There are many classic ways to remove outliers:

Trimming: Outlier Removal

There are many classic ways to remove outliers:

- ▶ Fit, remove outliers, refit
 - ▶ Upside: easy to do
 - ▶ Downside: outliers can affect the initial fit
 - ▶ Downside: when to stop?

Trimming: Outlier Removal

There are many classic ways to remove outliers:

- ▶ Fit, remove outliers, refit
 - ▶ Upside: easy to do
 - ▶ Downside: outliers can affect the initial fit
 - ▶ Downside: when to stop?
- ▶ Regression: replace least squares loss with 'robust' penalty
 - ▶ Upside: relatively easy to do
 - ▶ Downside: how to pick shape?
 - ▶ Downside: how to extend to non-additive errors?

Trimming: Outlier Removal

There are many classic ways to remove outliers:

- ▶ Fit, remove outliers, refit
 - ▶ Upside: easy to do
 - ▶ Downside: outliers can affect the initial fit
 - ▶ Downside: when to stop?
- ▶ Regression: replace least squares loss with 'robust' penalty
 - ▶ Upside: relatively easy to do
 - ▶ Downside: how to pick shape?
 - ▶ Downside: how to extend to non-additive errors?
- ▶ Our focus: trimming
 - ▶ Upside: works for any model
 - ▶ Upside: transparent assumptions
 - ▶ Downside: nonconvex model
 - ▶ Upside: doesn't seem to matter in practice

Trimming: Overview

Trimming uses auxiliary weights to detect outliers:

$$\min_{x, w} \sum_{i=1}^m w_i f_i(x) \quad \text{s.t.} \quad w_i \in [0, 1], \quad \sum_{i=1}^m w_i = h$$

- ▶ For fixed x , minimal h residuals have $w_i = 1$, rest are 0
- ▶ Minimal h residuals are thus classified as 'inliers'
- ▶ Remaining $m - h$ points are by default 'outliers'
- ▶ As x varies, we are looking to only fit inliers.

Problem is theoretically hard, but practically works very well.

Trimming: Overview


The general idea extends to any learning model:

$$\min_{x, w} \sum_{i=1}^m w_i f_i(x) \quad \text{s.t.} \quad w_i \in [0, 1], \quad \sum_{i=1}^m w_i = h$$

- ▶ Least squares: $f_i(x) = \frac{1}{2}(y_i - x_0 - x_1 a_i)^2$
- ▶ Logistic: $f_i(x) = \log(1 + \exp(x_0 + x_1 a_i)) - y_i(x_0 + x_1 a_i)$
- ▶ Neural net: $f_i(x) = \text{soft max for a labeled data point}$









Trimming: Least Squares Example

$$\min_{x,w} \sum_{i=1}^m \frac{w_i}{2} (y_i - x_0 - x_1 a_i)^2 \quad \text{s.t.} \quad w_i \in [0, 1], \quad \sum_{i=1}^m w_i = h$$









Trimming: CNN Example

Here we see the results of a convolutional neural network (CNN) classifier that predicts cats and birds, with inliers (top row) and outliers (bottom row) identified using trimming.



Further Reading I

-  Aravkin, Aleksandr and Damek Davis (2020). “Trimmed statistical estimation via variance reduction”. In: *Mathematics of Operations Research* 45.1, pp. 292–322.
-  Beck, Amir (2014). *Introduction to nonlinear optimization: Theory, algorithms, and applications with MATLAB*. SIAM.
-  Beck, Amir and Yonina C. Eldar (2013). “Sparsity constrained nonlinear optimization: Optimality conditions and algorithms”. In: *SIAM Journal on Optimization* 23.3, pp. 1480–1509.
-  Beck, Amir and Marc Teboulle (2009). “A fast iterative shrinkage-thresholding algorithm for linear inverse problems”. In: *SIAM Journal on Imaging Sciences* 2.1, pp. 183–202.
-  Boyd, Stephen and Lieven Vandenberghe (2004). *Convex optimization*. Cambridge University Press.
-  — (2018). *Introduction to applied linear algebra: Vectors, matrices, and least squares*. Cambridge University Press.

Further Reading II



Candes, Emmanuel J., Justin K. Romberg, and Terence Tao (2006). “Stable signal recovery from incomplete and inaccurate measurements”. In: *Communications on Pure and Applied Mathematics: A Journal Issued by the Courant Institute of Mathematical Sciences* 59.8, pp. 1207–1223.



Hosmer Jr., David W., Stanley Lemeshow, and Rodney X. Sturdivant (2013). *Applied logistic regression*. Vol. 398. John Wiley & Sons.



Kelley, Carl T. (1999). *Iterative methods for optimization*. SIAM.



Mairal, Julien et al. (2009). “Supervised dictionary learning”. In: *Advances in Neural Information Processing Systems*, pp. 1033–1040.

Further Reading III



Neykov, Neyko M. and Christine H. Müller (2003). “Breakdown point and computation of trimmed likelihood estimators in generalized linear models”. In: *Developments in Robust Statistics*. Springer, pp. 277–286.



Nocedal, Jorge and Stephen Wright (2006). *Numerical optimization*. Springer Science & Business Media.



Parikh, Neal and Stephen Boyd (2014). “Proximal algorithms”. In: *Foundations and Trends in Optimization* 1.3, pp. 127–239.



Rousseeuw, Peter J. (1985). “Multivariate estimation with high breakdown point”. In: *Mathematical Statistics and Applications* 8.283-297, p. 37.



Tibshirani, Robert (2011). “Regression shrinkage and selection via the lasso: a retrospective”. In: *Journal of the Royal Statistical Society: Series B (Statistical Methodology)* 73.3, pp. 273–282.

Further Reading IV



Tropp, Joel A. and Stephen J. Wright (2010). “Computational methods for sparse solution of linear inverse problems”. In: *Proceedings of the IEEE* 98.6, pp. 948–958.



Yang, Eunho, Aurélie C Lozano, Aleksandr Aravkin, et al. (2018). “A general family of trimmed estimators for robust high-dimensional data analysis”. In: *Electronic Journal of Statistics* 12.2, pp. 3519–3553.